

BERT

Pre-training of Deep Bidirectional Transformers for
Language Understanding

Overview

- Define what are pre-trained language representation models
- Understand BERT architecture
- Which is Based on Transformers, which includes :
 - Which includes Encoder - Decoder
 - Attention and Self-Attention
-
- And finally, we'll see an application of BERT for document classification

Pre-trained language representation models

Language representation... We've got word2vec already!

- True ! But ...
- Unable to process unknown or out-of-vocabulary (OOV) words.
- Not a multilingual models (requires new embedding matrices and not allow for parameters sharing)
- Represents every word as an independent vector
 - Only captures weak relations between words
 - No difference with “bank account” and “the bank of the river”
- We need to build more complex relationship than word encoding

Why pre-trained ?

- Many NLP tasks lack labeled data specific to the task.
- It's an issue as deep learning-based models benefit from training over millions or even billions of annotated examples.
- **Global idea over the past few years : train general purpose language representation models.**
- Training in two phase : **pre-training** over enormous amount of unannotated data from texts and web and **fine-tune** it on a smaller dataset corresponding to the downstream task.

Example : sentence classification (Spam or not Spam)

1 - **Semi-supervised** training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.

Semi-supervised Learning Step

Model:



Dataset:



Objective:

Predict the masked word
(language modeling)

2 - **Supervised** training on a specific task with a labeled dataset.

Supervised Learning Step

Classifier

75% Spam
25% Not Spam

Model:
(pre-trained
in step #1)



Dataset:

Email message	Class
Buy these pills	Spam
Win cash prizes	Spam
Dear Mr. Atreides, please find attached...	Not Spam

BERT's architecture

BERT : Bidirectional Encoder Representations from Transformers


- BERT is made of two big stacks : one of Encoders and one of Decoders
- Actually its architecture is based on the **Transformer**'s architecture
- The
- So, let's start by understanding what a **Transformer** is...

The Transformer

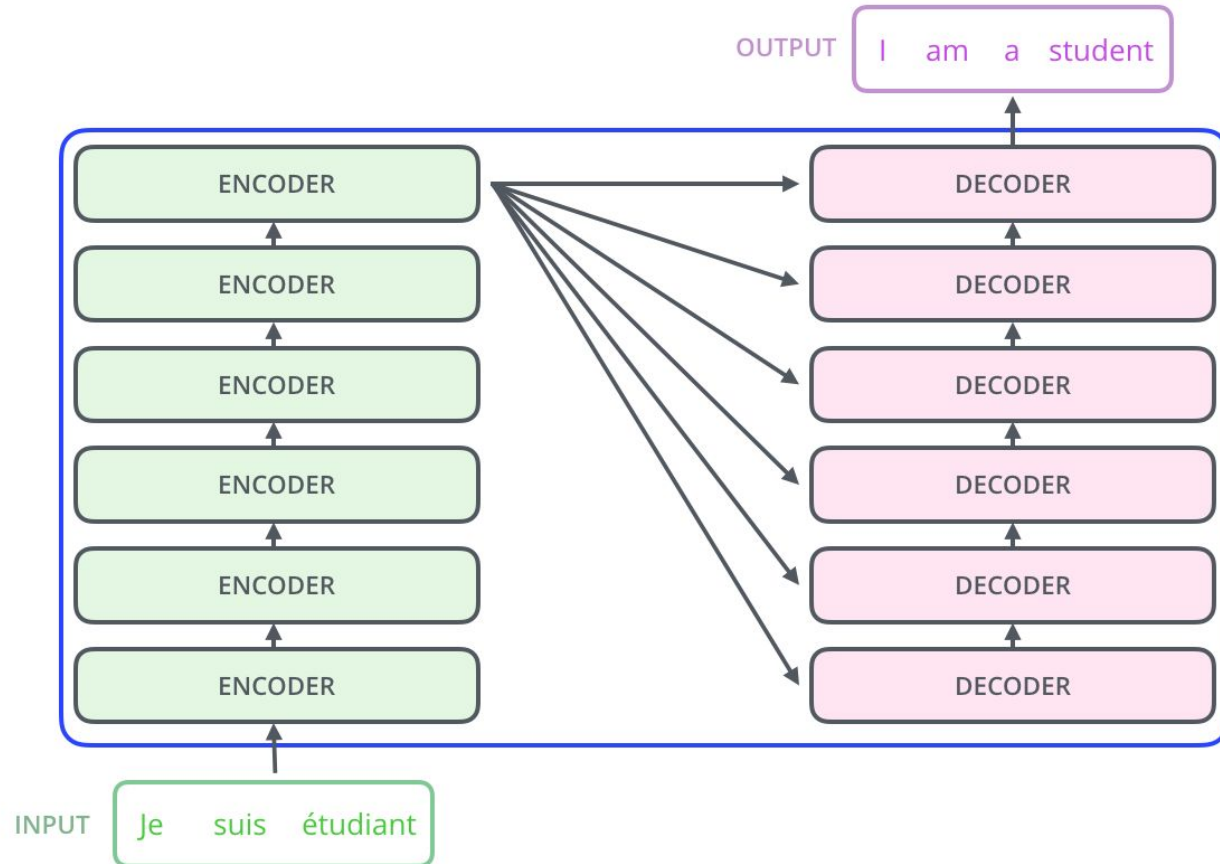
Transformer - Overview as a black box



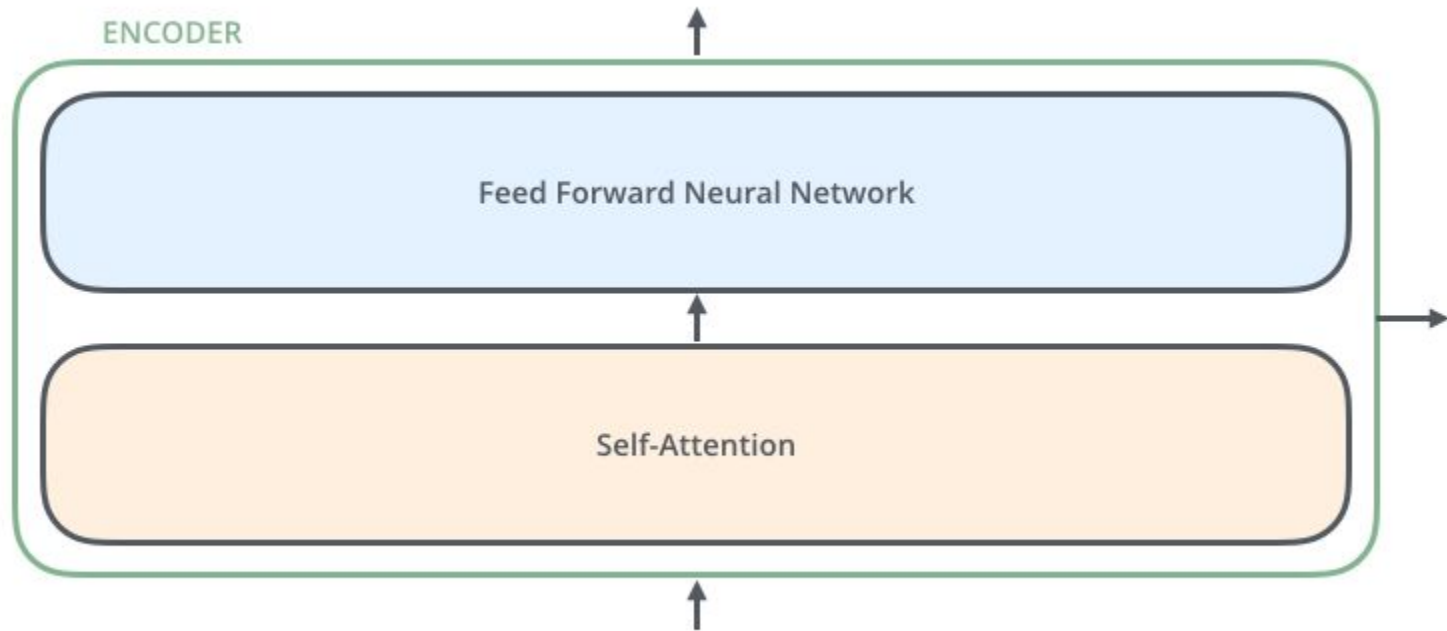
Transformer - Encoding component and Decoding component



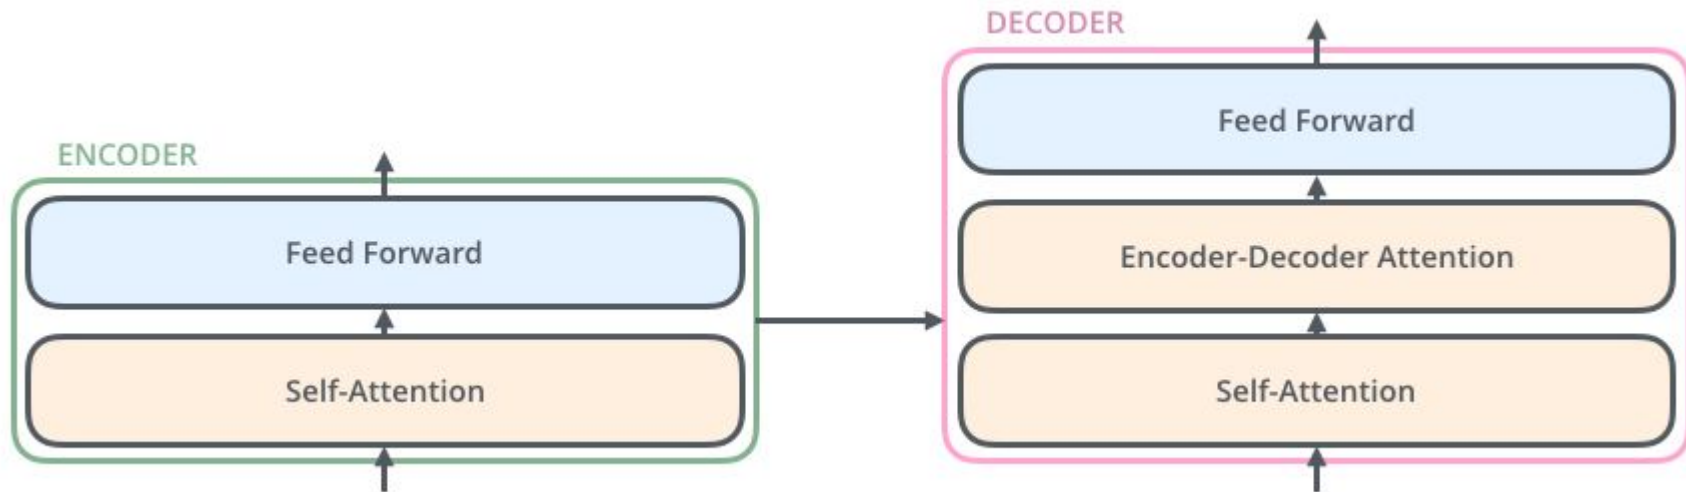
Transformer - Encoding component and Decoding component



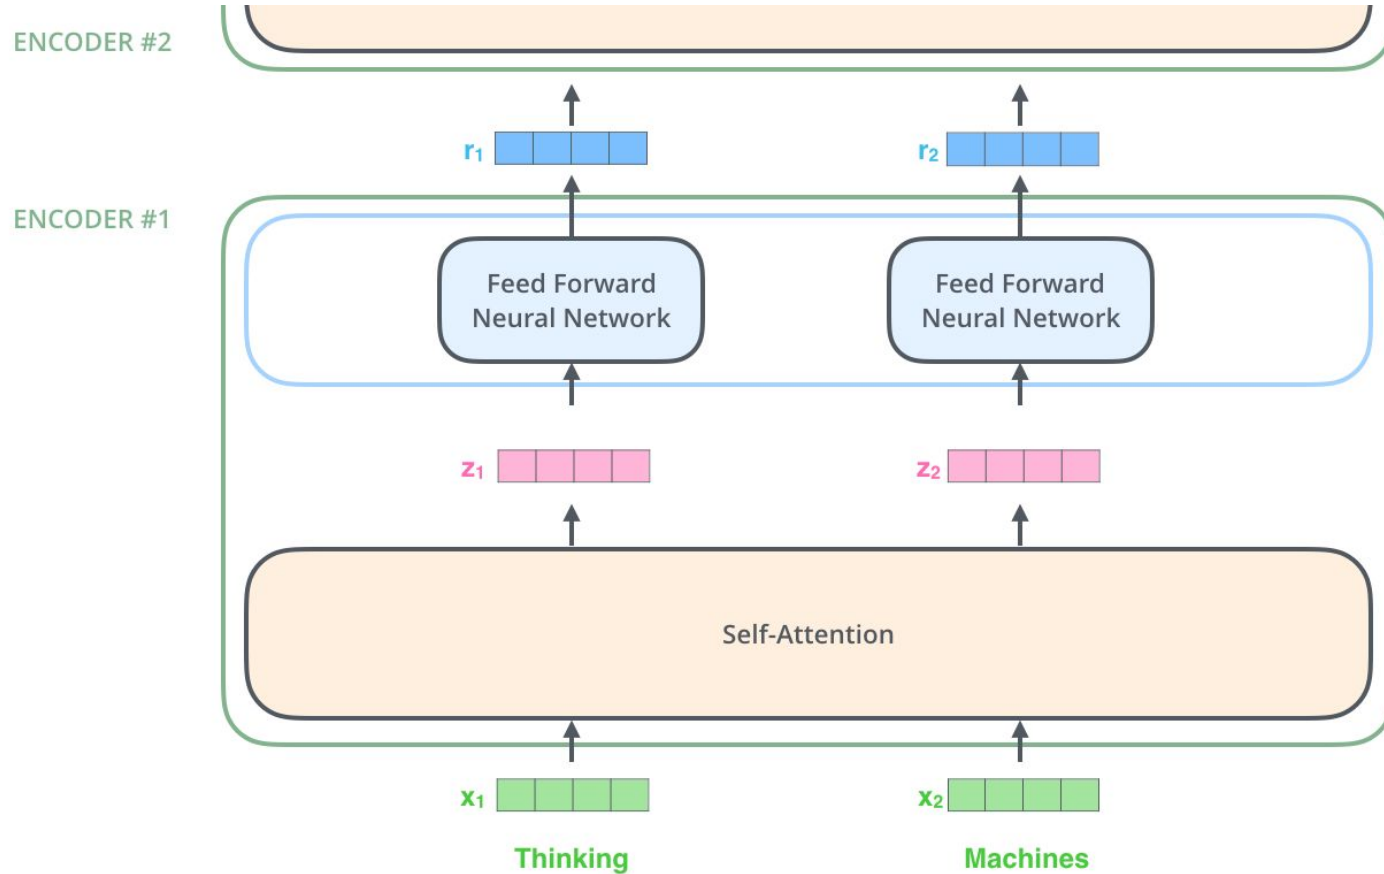
Transformer - Encoder



Transformer - Encoder and Decoder

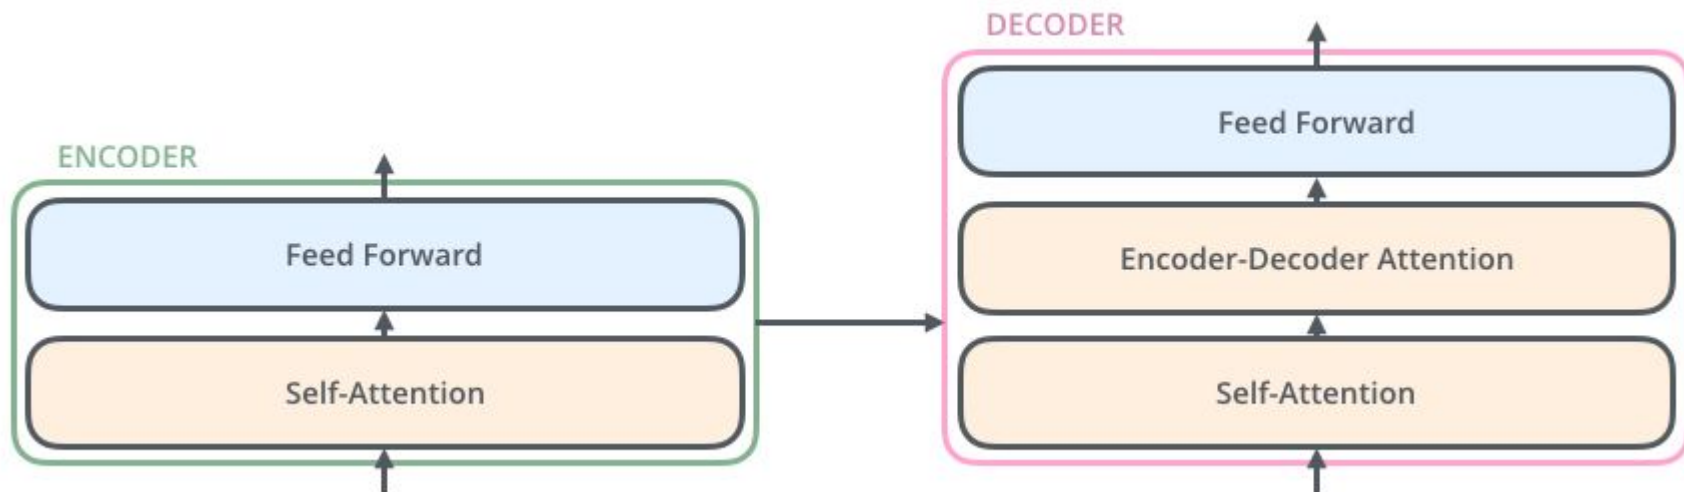


Transformer - Inputs and Outputs

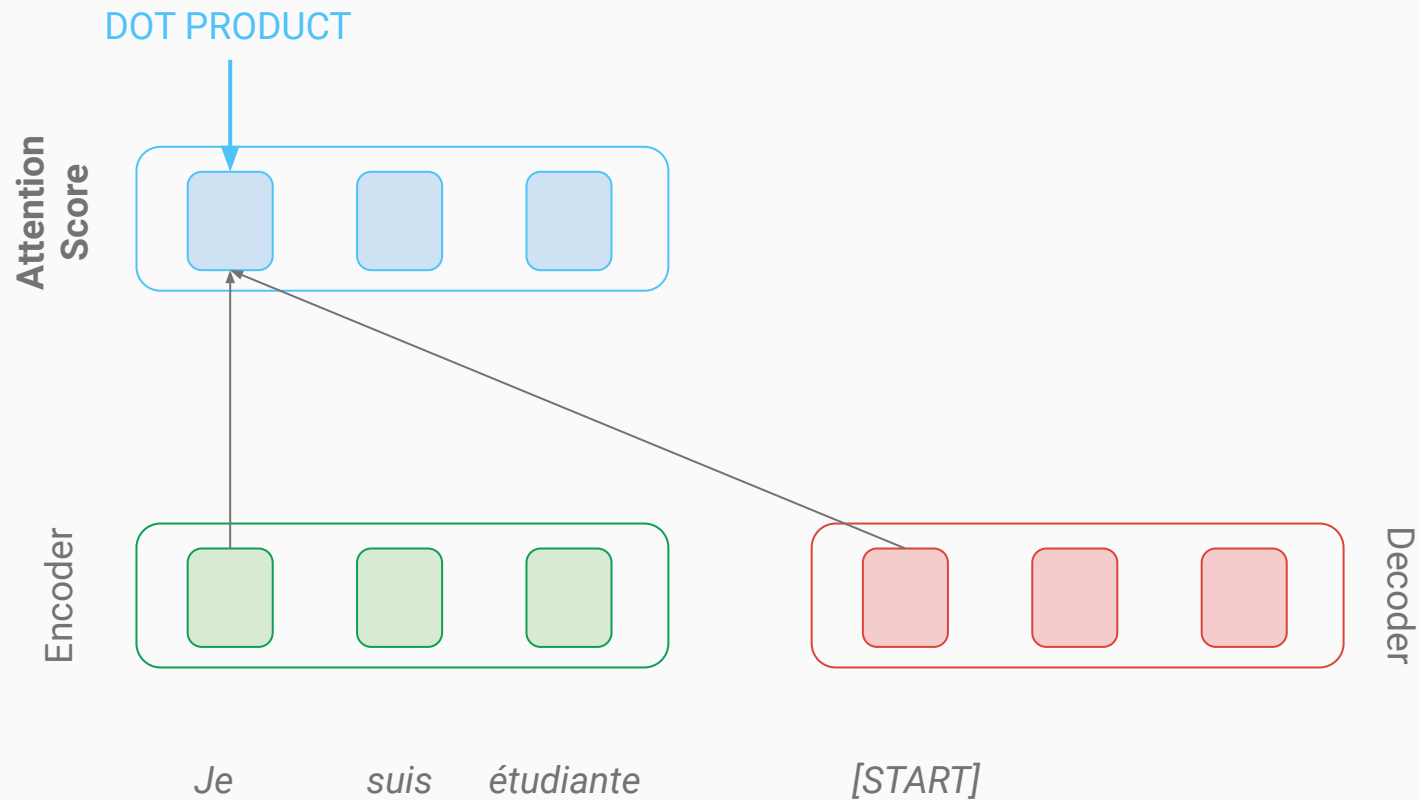


Attention and self-Attention

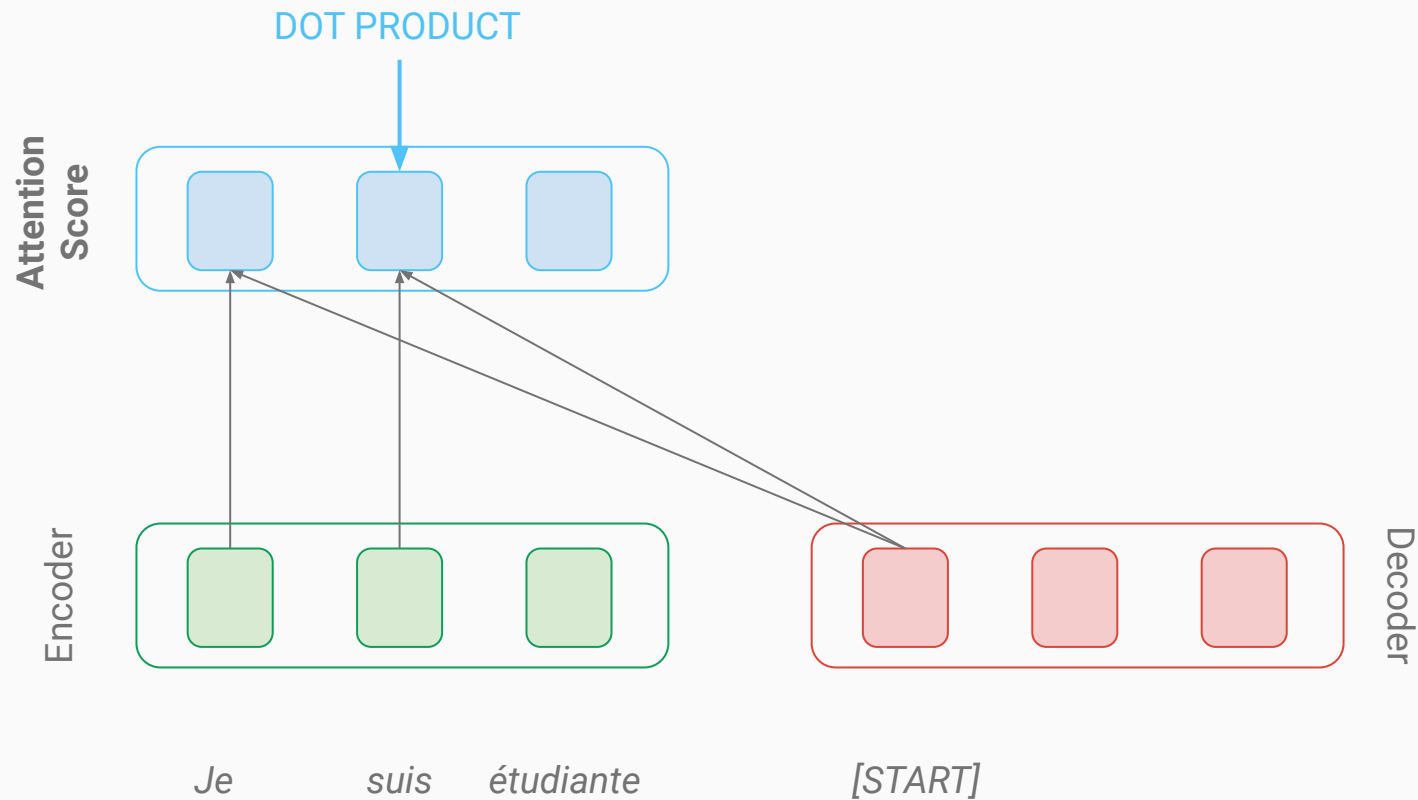
Transformer - Remember, Attention and self-Attention



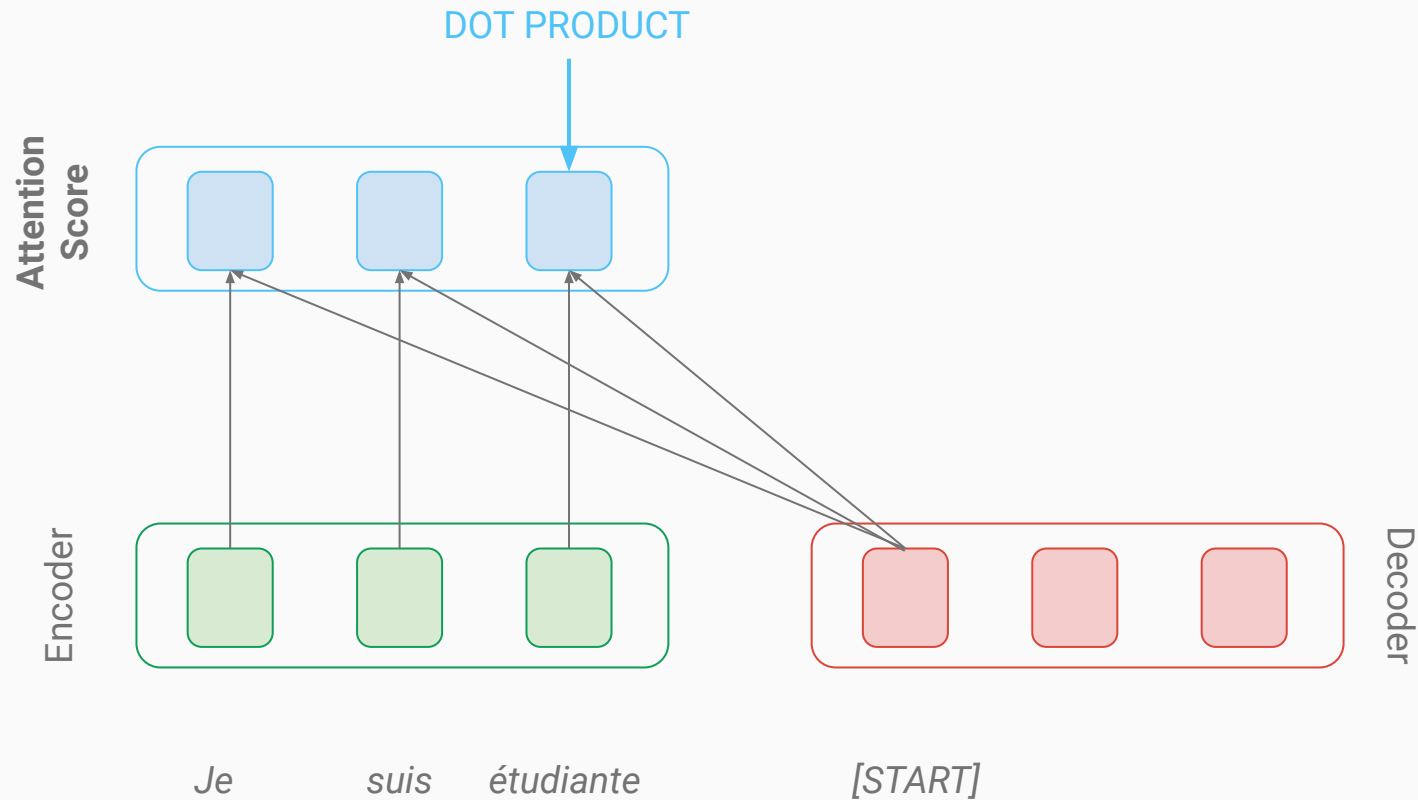
Transformer - Attention



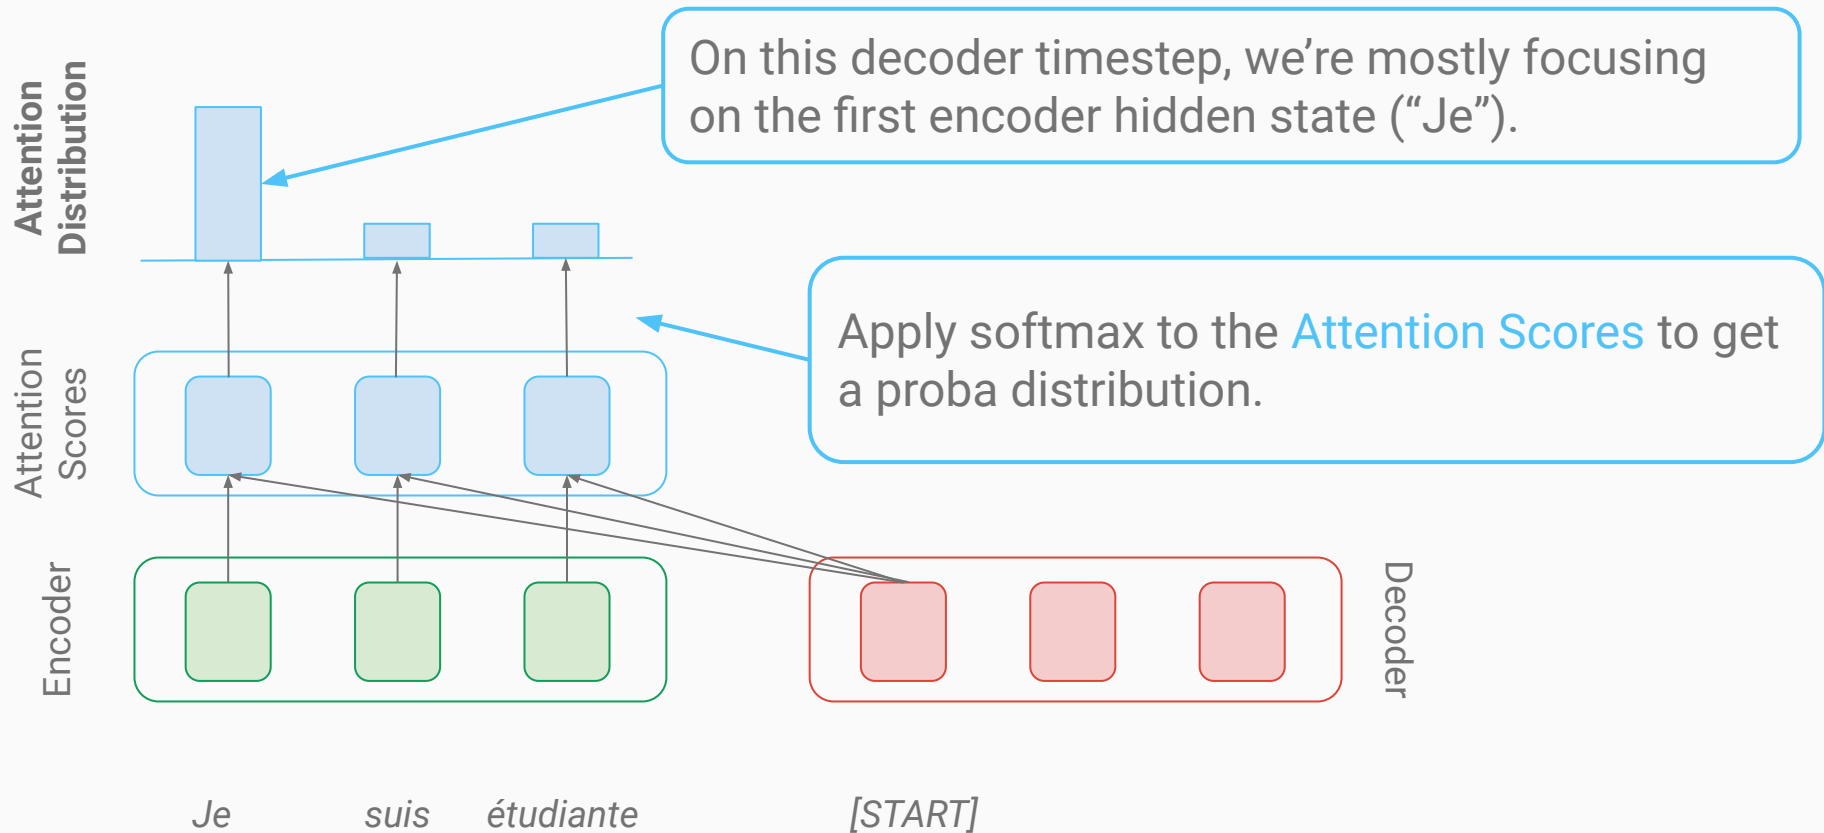
Transformer - Attention



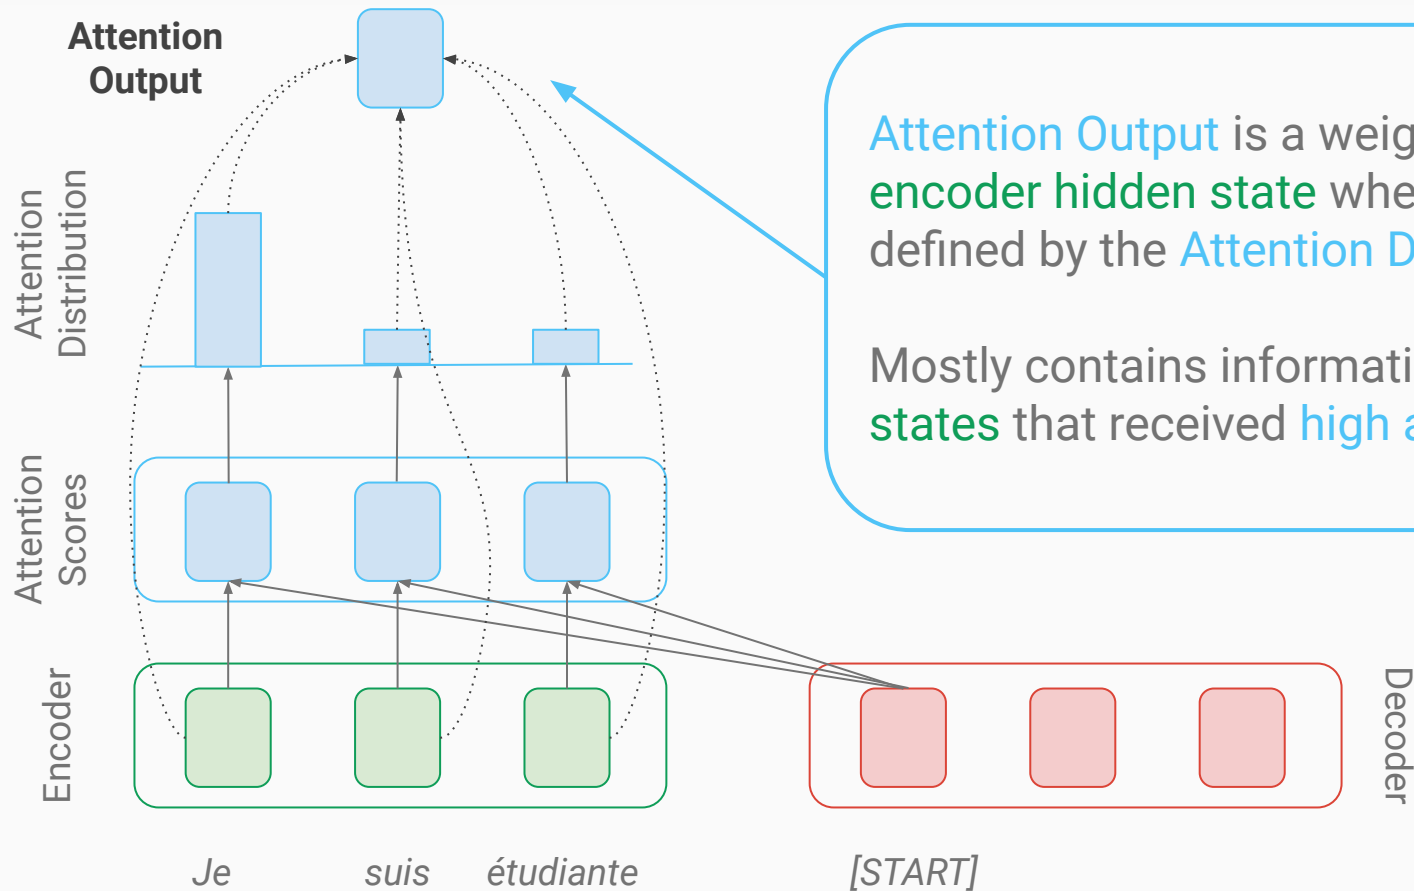
Transformer - Attention



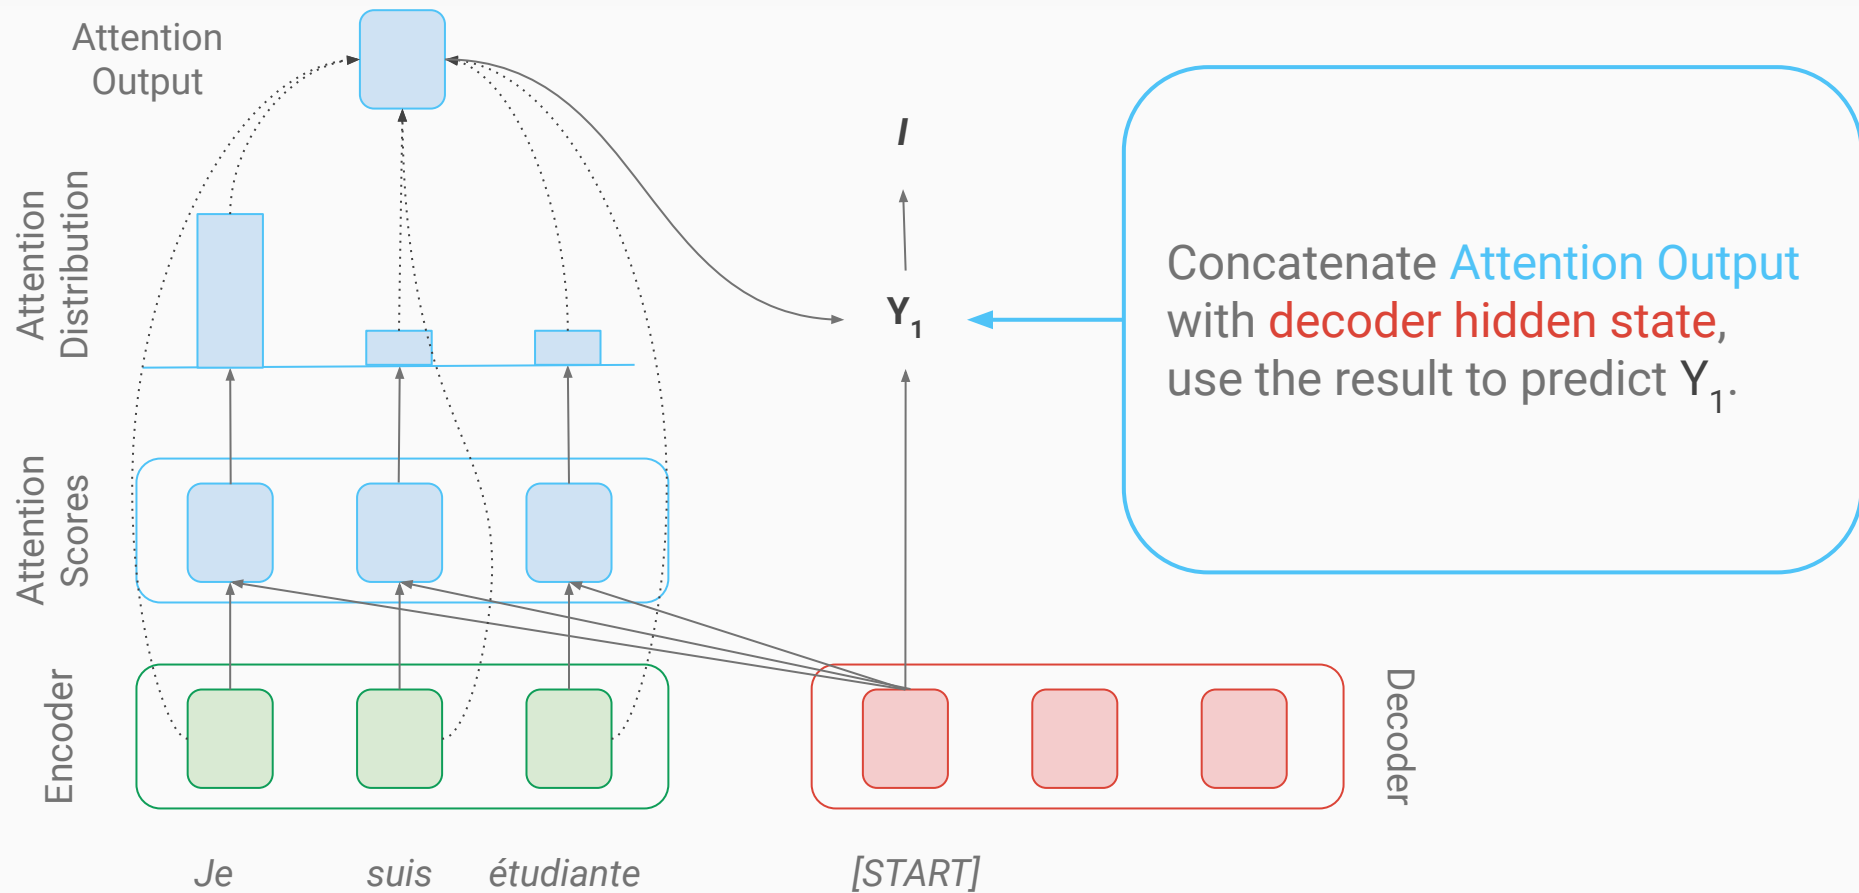
Transformer - Attention



Transformer - Attention



Transformer - Attention



Generalize Attention Definition

Given a set of **vector values**, and a **vector query**, **Attention** is a technique to compute a weighted sum of the values, dependent on the query.

In our case, we had decoder hidden state attending to encoder hidden state :

- Queries -> **Decoder hidden state**
- Values -> **Encoder hidden state**

Now, let's see how we it works with self-Attention.

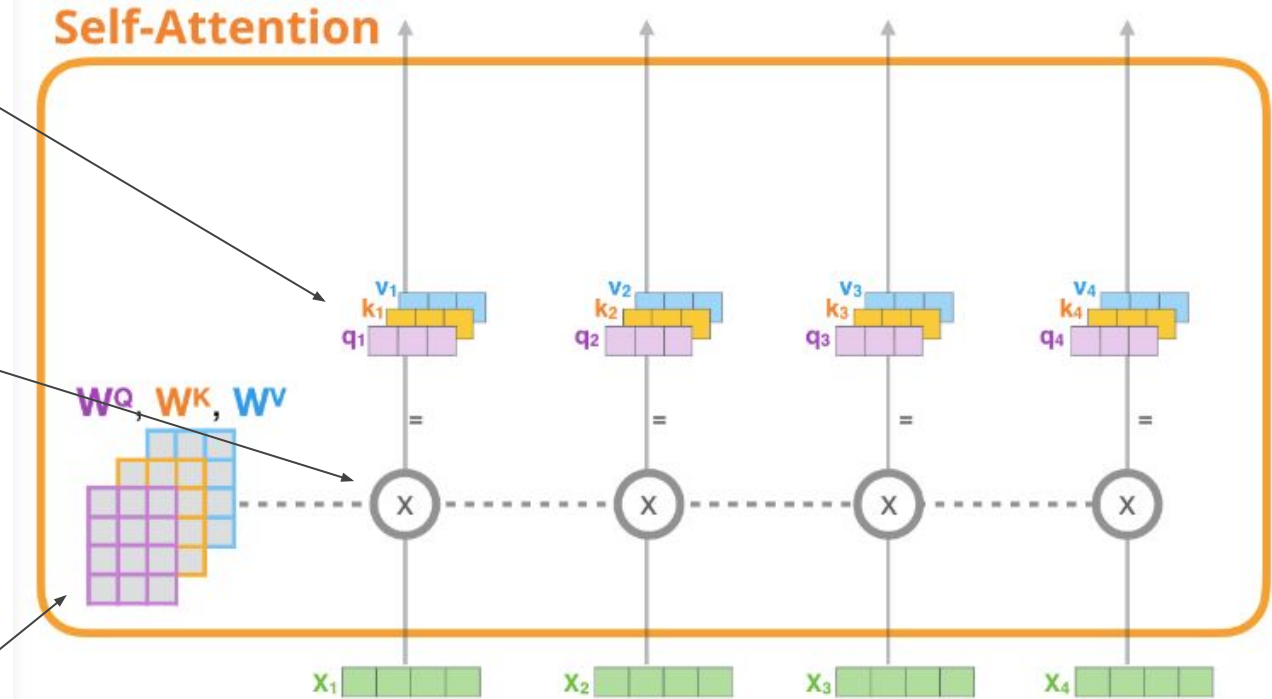
Transformer - Self-Attention : Queries, Keys and Values

Which builds our
Query vectors, Key
vectors and Value
vectors

Multiply each input
vectors (word
embedded) by each
matrices.

Three matrices (init
random)

W^Q , W^K , W^V

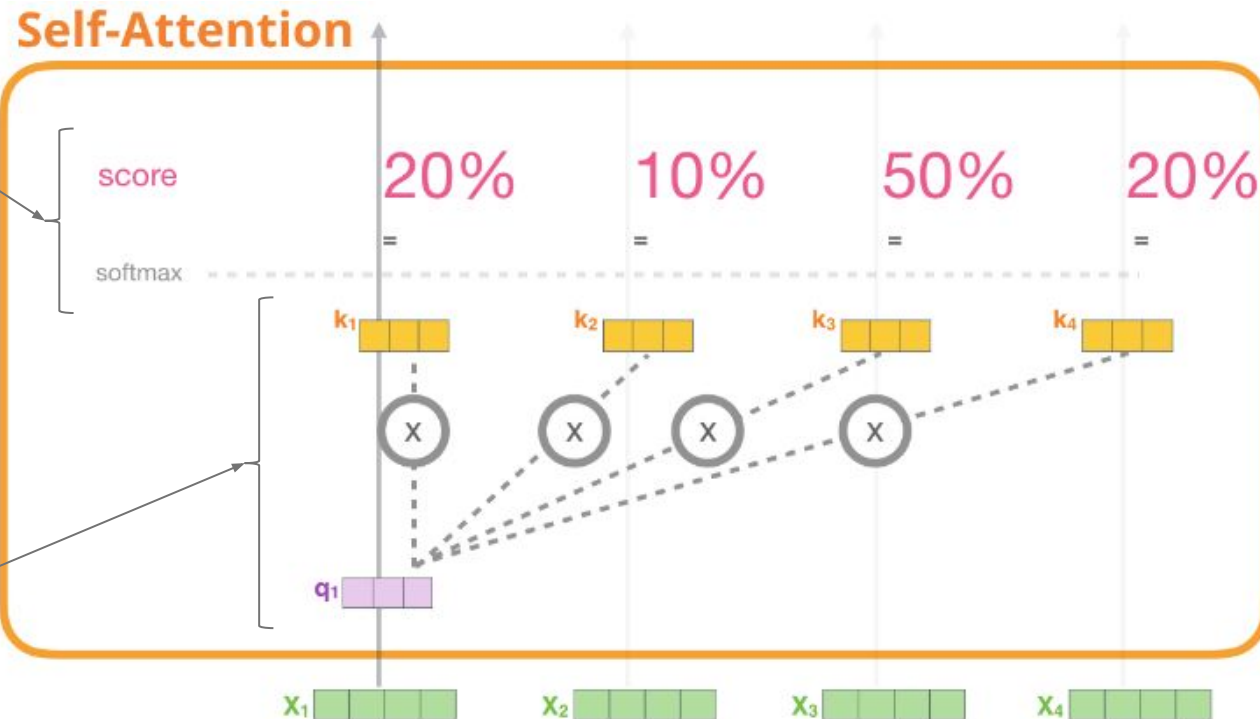


Transformer - Self-Attention : Compute Attention Score and Attention Distribution

Apply softmax on all attention scores to get the **Attention Distribution**.

$$\text{Attention score} = q_i \cdot k_i$$

Scores are normalised afterward by the square root of the dimension of the key vectors

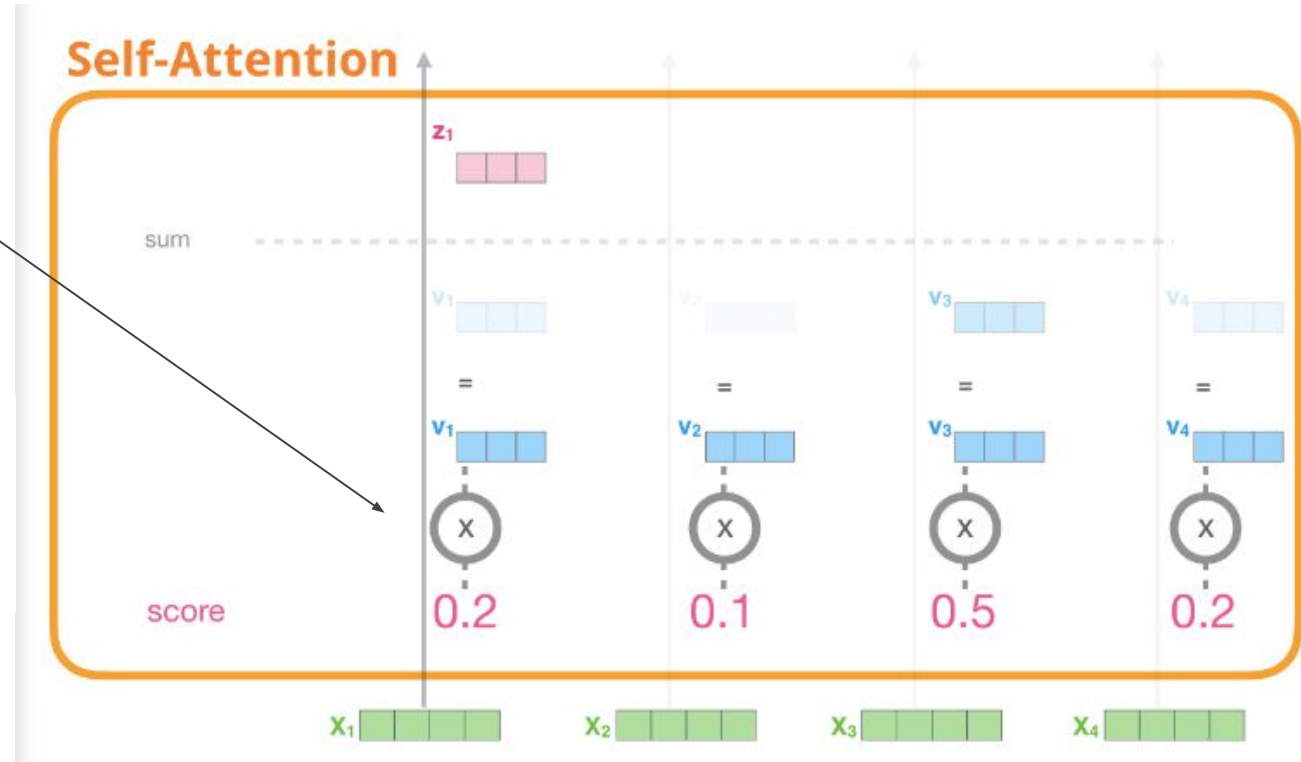


Transformer - Self-Attention : Compute Attention Output

Attention output = value
vector x attention
distribution

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) \times V = Z$$

Self-Attention




Multi-Head Self-Attention

Transformer - Multi-Head Self-Attention

Gives attention layer multiple “representation subspaces” by building multiple sets of **Query**, **Key**, **Value** weight matrices

- The transformer uses 8 attention head
- So each encoder and decoder as 8 set of matrices.

Each sets are randomly initialized and then trained.




Transformer - Multi-Head Self-Attention

1) Concatenate all the attention heads



2) Multiply with a weight matrix W^O that was trained jointly with the model

\times



3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



Transformer - Multi-Head Self-Attention : recap


1) This is our input sentence*

Thinking
Machines


2) We embed each word*





3) Split into 8 heads.
We multiply X or R with weight matrices



4) Calculate attention using the resulting $Q/K/V$ matrices





5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



...

...

...



W^O

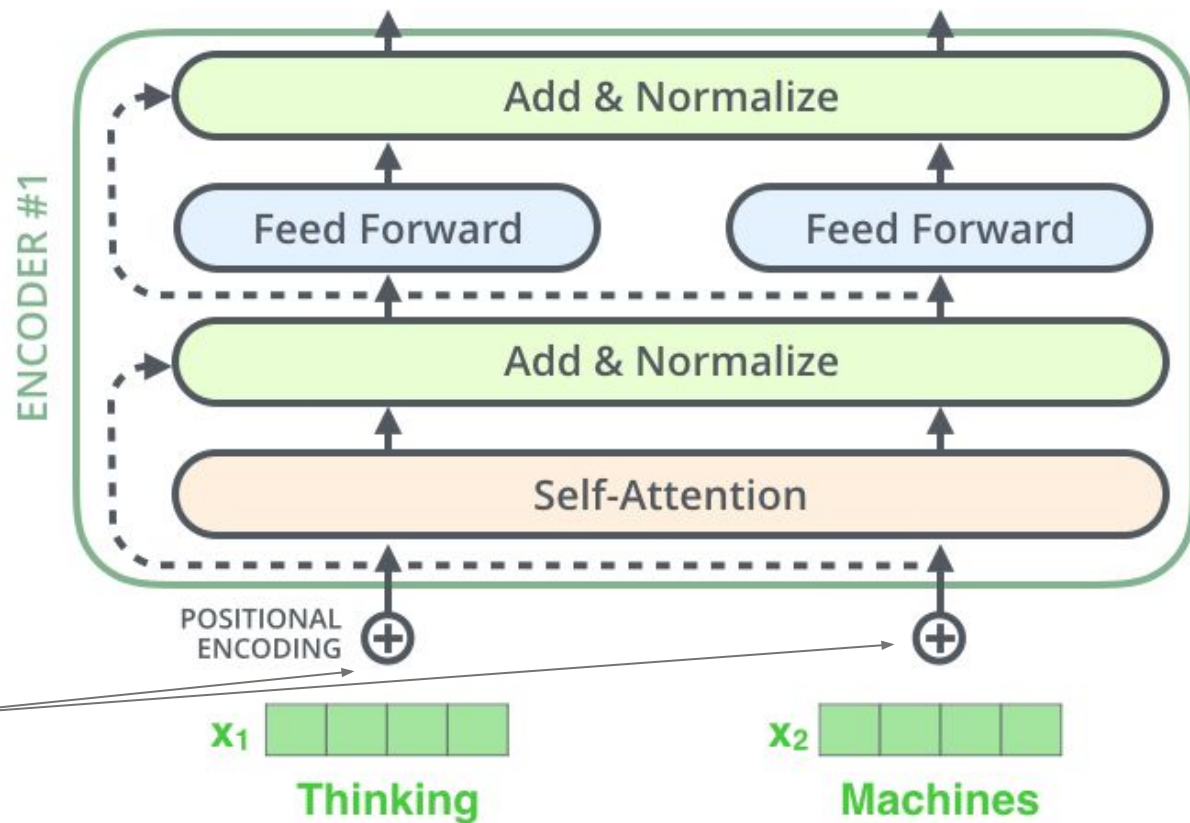


* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



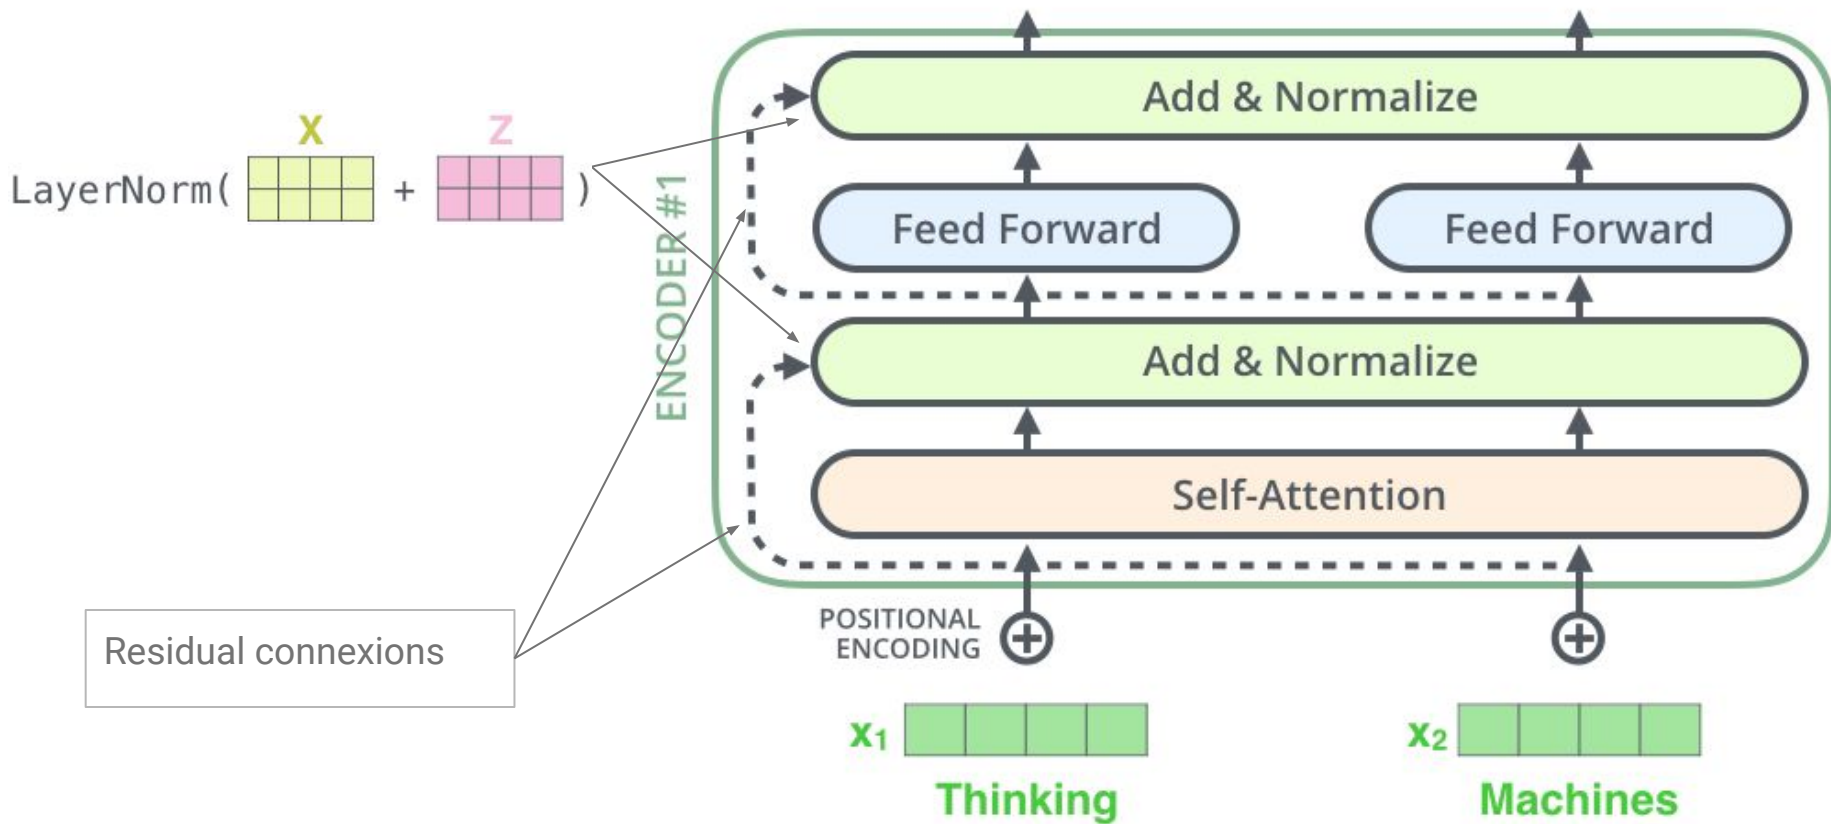
Transformer : almost done

Transformer - Positional encoding

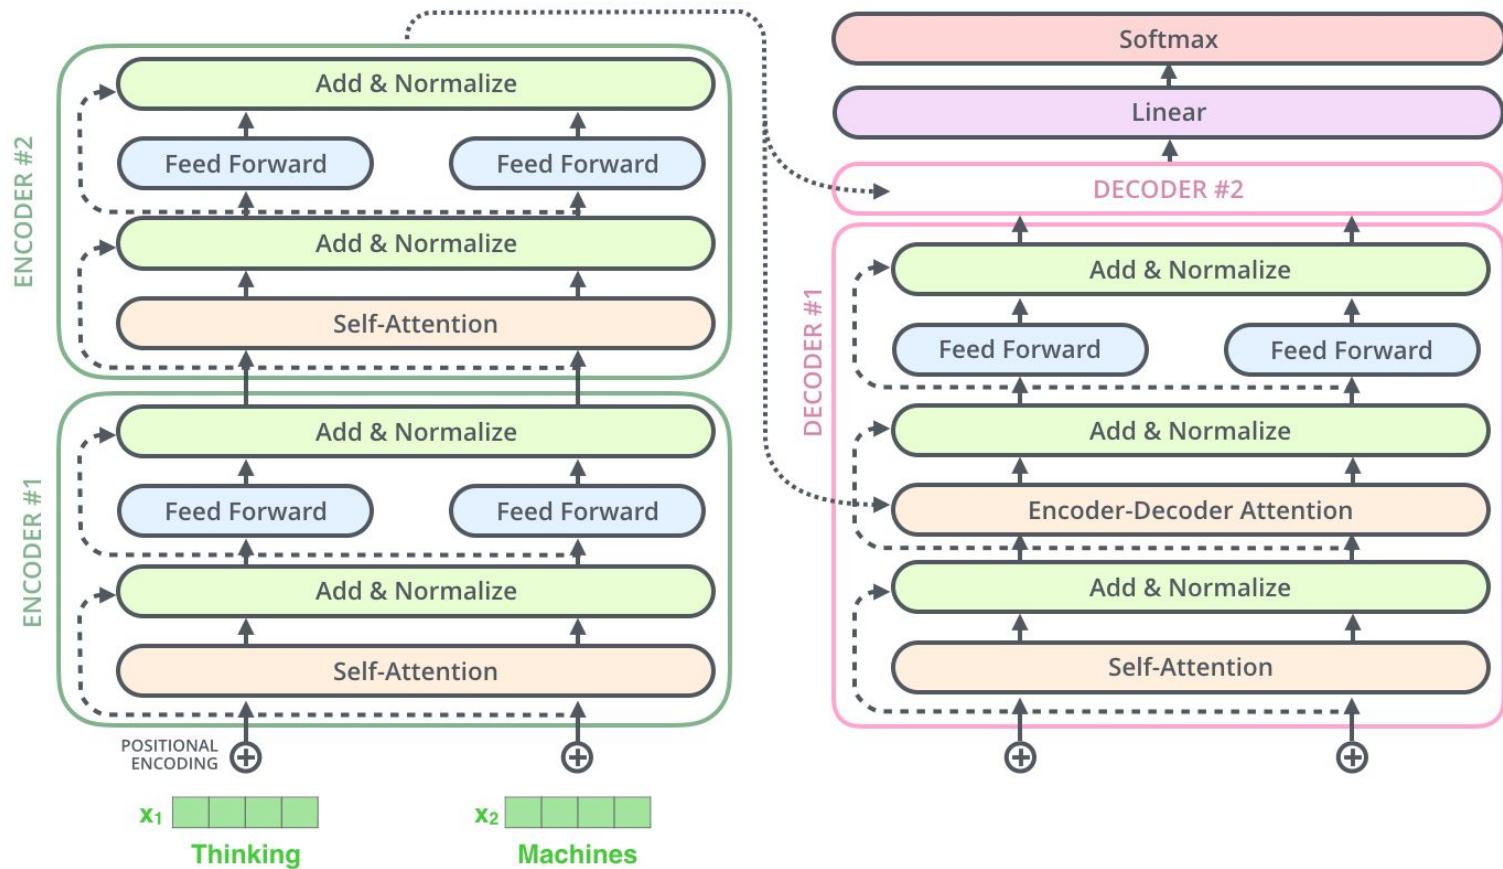


Add to each **Input vectors** (word embedded) a **positional encoding**

Transformer - Residual connexions for the “Add & Normalize” Layer



Transformer - Workflow



Transformer - Key ideas to keep in mind

- The **Transformer** is **build** with Encoders and Decoders, without any RNN nor CNN but **solely with Attention**.

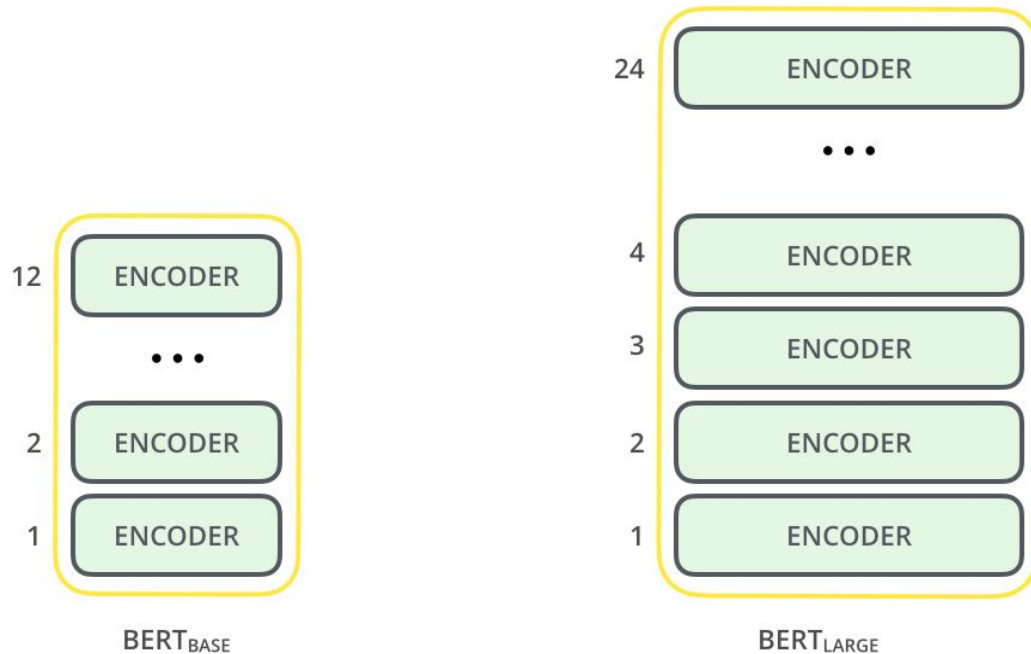
Multi-Head Self-Attention + Positional Encoding

= Context influence + Order in sequence = Sequence Representation

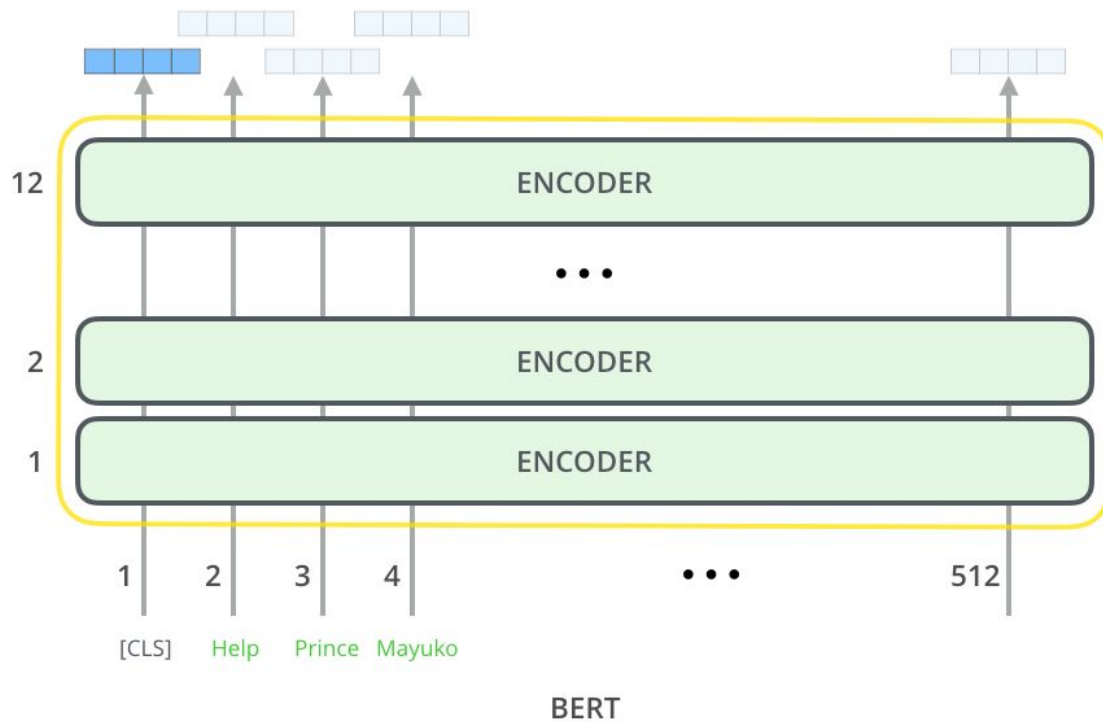
- Attention mechanism solve the **bottleneck problem** by allowing the decoder to look directly at the source.

Back to BERT

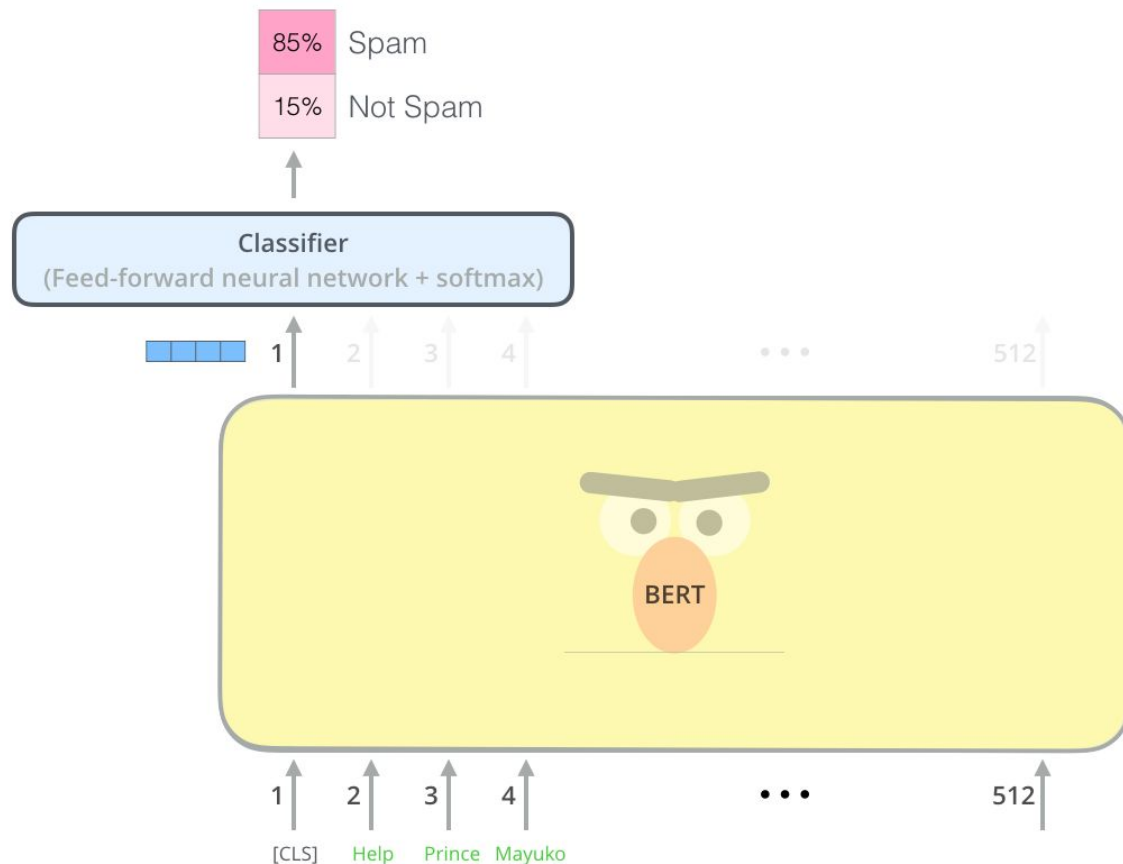
BERT architecture - Big Trained Transformer Encoder Stack



BERT - Input & Output



BERT - Fine-Tuning layer for classification (Spam or Not Spam)



Bert - Pre-training a deep bi-directional language representation

Masked LM:

Each sentences, 15% words replace by [MASK]

problem: In reality the [MASK] doesn't Exist:

Solution: the N% selected words

- 80% of time become the Token [MASK]
- 10% replaced by other Token
- 10% Token not change

PPS:

Sentence pair combination A&B.

In the Training set, 50% are good combinaison, other Random Combination of same corpus

Inject 2 Tokens

- <CLS> Start of A
- <SEP> End each sentence

Compute the probability that A&B is good combinaison.





"That's all Folks!"

Questions ?

Further readings

- Fine-Tuning

- [Semi-supervised Sequence Learning](#) : Paper introducing improvement in sequence learning with RNN

- Transformers

- [Attention Is All You Need](#) : Paper introducing the Transformer
- [The Annotated Transformer](#) : Step by step annotated in python of “Attention in all you need”

- BERT

- [Open Sourcing BERT: State-of-the-Art Pre-training for Natural Language Processing](#) : Google blog post announcing BERT release
- [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) : The paper
- [The Illustrated BERT, ELMo, and co. \(How NLP Cracked Transfer Learning\)](#) : A good blogpost ;)